

# Forecasting Bitcoin Price Using ARIMA Model

*Longyin Poon*

*4/10/2018*

**Training dataset includes 230 weeks from 4/8/2013 to 8/21/2017**

**Testing dataset consists 30 weeks of data from 8/28/2017 to 3/26/2018**


```
# Load dataset
bitcoin_data <- read_csv("bitcoin_data.csv")
training <- bitcoin_data[1:230, ]
testing <- bitcoin_data[231:260, ]
```

## Time Series Analysis

**Stationary Test**

```
# Stationary Test
tsplot(x = training$week, y = training$bitcoin_price, ylab = "Bitcoin Price",
       xlab = "Number of Weeks", main = "Weekly Bitcoin Price from 4/8/2013 to 8/21/2017",
       col = "steelblue", lwd = 3)
```

## Weekly Bitcoin Price from 4/8/2013 to 8/21/2017



```
adf.test(bitcoin_ts, alternative = "stationary")
```

```
## Warning in adf.test(bitcoin_ts, alternative = "stationary"): p-value  
## greater than printed p-value
```


```
##  
## Augmented Dickey-Fuller Test  
##  
## data: bitcoin_ts  
## Dickey-Fuller = 3.5991, Lag order = 6, p-value = 0.99  
## alternative hypothesis: stationary
```

```
# Fail to reject the series is non-stationary with high  
# p-value
```

### Log transformation

```
training$log_price <- log(training$bitcoin_price)  
tsplot(x = training$week, y = training$log_price, ylab = "Logged Bitcoin Price",  
       xlab = "Number of Weeks", main = "Weekly Logged Bitcoin Price from 4/8/2013 to 8/21/2017",  
       col = "steelblue", lwd = 3)
```

## Weekly Logged Bitcoin Price from 4/8/2013 to 8/21/2017




```
# The augmented Dickey-Fuller test
adf.test(training$log_price, alternative = "stationary")

##
## Augmented Dickey-Fuller Test
##
## data: training$log_price
## Dickey-Fuller = -1.3664, Lag order = 6, p-value = 0.8421
## alternative hypothesis: stationary
# Fail to reject the series is non-stationary but p-value is
# lower
```

## ACF and PACF for logged Bitcoin Price


```
Acf(training$log_price, main = "ACF for Logged Bitcoin Price")
```

### ACF for Logged Bitcoin Price



```
# Significant autocorrelations with many lags caused by carry  
# over from earlier lags  
Pacf(training$log_price, main = "PACF for Logged Bitcoin Price")
```

### PACF for Logged Bitcoin Price




```
# A spike at 1 in PACF meaning lag 1 alone explained most of  
# the information
```

First order differencing on logged Bitcoin Price

```
diff_log_price <- diff(training$log_price)  
plot(diff_log_price, type = "l", xlab = "Number of Weeks", ylab = "Differenced logged Bitcoin Price",  
     main = "First Order Differencing on logged Bitcoin Price")
```

### First Order Differencing on logged Bitcoin Price



```
adf.test(diff_log_price, alternative = "stationary")
```

```
## Warning in adf.test(diff_log_price, alternative = "stationary"): p-value  
## smaller than printed p-value
```


```
##  
## Augmented Dickey-Fuller Test  
##  
## data: diff_log_price  
## Dickey-Fuller = -4.9599, Lag order = 6, p-value = 0.01  
## alternative hypothesis: stationary
```

```
# Reject H0 with low p-value. The series is stationary
```

ACF and PACF for Differenced log Bitcoin Price


```
# ACF and PACF for differenced series to determine p and q  
# for ARIMA model  
Acf(diff_log_price, main = "ACF for Differenced log Price")
```

### ACF for Differenced log Price



```
# Significant spike at lag 1 and lag 2. Consider MA(2) model
Pacf(diff_log_price, main = "PACF for Differenced log Price")
```

### PACF for Differenced log Price

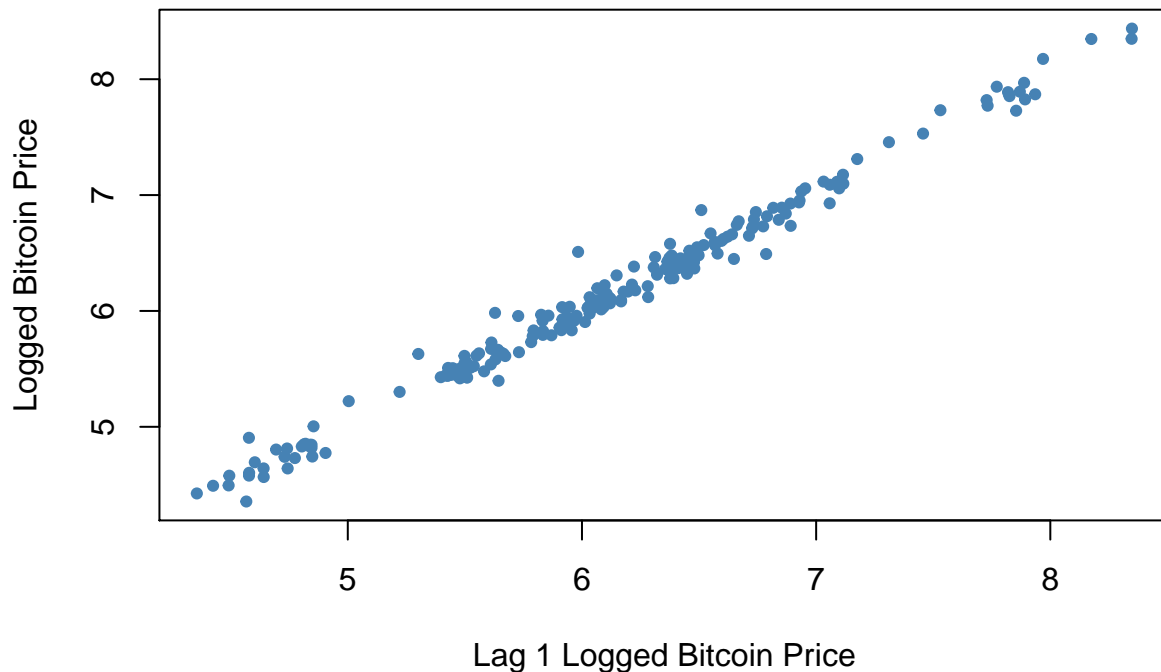


```
# Significant spike at lag 1 and close to significant for lag
# 2. Consider AR(1) or AR(2) model
```

Scatterplot of logged Bitcoin Price against its lags

```
plot(x = lag(training$log_price[2:230]), y = training$log_price[2:230],  
     ylab = "Logged Bitcoin Price", xlab = "Lag 1 Logged Bitcoin Price",  
     pch = 19, cex = 0.7, col = "steelblue", main = "Scatterplot of Logged Bitcoin Price VS Lag 1 Logged Bitcoin Price")
```

## Scatterplot of Logged Bitcoin Price VS Lag 1 Logged Bitcoin Price




```
# Linear pattern suggesting the first order autoregressive  
# model is appropriate.
```

## ARIMA Model

```
## Series: training$log_price  
## ARIMA(2,1,1) with drift  
##  
## Coefficients:  
##          ar1      ar2      ma1      drift  
##        -0.3461  0.3146  0.6548  0.0142  
## s.e.    0.4411  0.1298  0.4607  0.0099  
##  
## sigma^2 estimated as 0.008916:  log likelihood=217.43  
## AIC=-424.85   AICc=-424.59   BIC=-407.69
```


## (2,1,0) Model Residuals



```
## Series: training$log_price
## ARIMA(1,1,2)
##
## Coefficients:
##      ar1      ma1      ma2
##    0.1657  0.1532  0.1692
## s.e.  0.4284  0.4225  0.1589
##
## sigma^2 estimated as 0.008973:  log likelihood=216.19
## AIC=-424.38   AICc=-424.2   BIC=-410.64
```




### (1,1,2) Model Residuals



Check Residuals for model validity


```
res <- fit$residuals
Acf(res, main = "ACF for Residuals")
```

### ACF for Residuals



```
Pacf(res, main = "PACF for Residuals")
```

### PACF for Residuals



```
# Based on ACF and PACF, residuals are randomly distributed.
```

```
Box.test(res, type = "Ljung", lag = 1)
```

```
##
## Box-Ljung test
##
## data: res
## X-squared = 0.093094, df = 1, p-value = 0.7603
# P-value is high so we fail to reject H0: The residuals are
# independent. Box-Ljung test confirms our results.
```


Antilog function to reverse log transformation

```
# Anti log function
antilog <- function(lx, base) {
  lbx <- lx/log(exp(1), base = base)
  result <- exp(lbx)
  result
}
```

Forecast

```
fcast <- forecast(fit, h = 30)
plot(fcast)
```


### Forecasts from ARIMA(2,1,1) with drift



```
# Take the exponent
fcast$mean <- antilog(fcast$mean)
fcast$lower <- antilog(fcast$lower)
fcast$upper <- antilog(fcast$upper)
fcast$x <- antilog(fcast$x)
```

```
plot(fcast)
```

## Forecasts from ARIMA(2,1,1) with drift



Compare prediction and actual bitcoin price for 30 weeks

```
actual <- testing$bitcoin_price
pred <- fcast$mean[1:30]
error <- abs(actual - pred)/actual
prediction_table <- cbind(actual, pred, error)
colnames(prediction_table) <- c("Actual", "Prediction", "%Change")
prediction_table <- as.data.frame(prediction_table)
prediction_table$Week <- seq(1:30)
prediction_table
```


##	Actual	Prediction	%Change	Week
## 1	4448.186	4749.975	0.06784529	1
## 2	3877.361	4905.766	0.26523342	2
## 3	3836.227	4968.332	0.29510882	3
## 4	4104.402	5070.842	0.23546417	4
## 5	4349.547	5129.876	0.17940472	5
## 6	5163.845	5218.222	0.01053046	6
## 7	5724.994	5283.266	0.07715767	7
## 8	5764.160	5367.087	0.06888661	8
## 9	6814.830	5437.888	0.20205077	9
## 10	6936.828	5520.471	0.20417927	10
## 11	7254.876	5595.853	0.22867686	11
## 12	8305.702	5678.743	0.31628382	12
## 13	10326.294	5757.848	0.44240904	13
## 14	13732.345	5841.911	0.57458749	14
## 15	17376.360	5924.226	0.65906407	15
## 16	16223.990	6009.992	0.62956138	16

```
## 17 14485.893    6095.232 0.57922979    17
## 18 15325.235    6183.044 0.59654490    18
## 19 14512.732    6271.070 0.56789182    19
## 20 12177.336    6361.160 0.47762307    20
## 21 11282.490    6451.919 0.42814763    21
## 22  9788.082    6544.454 0.33138545    22
## 23  8014.081    6637.945 0.17171468    23
## 24  9670.893    6733.058 0.30378110    24
## 25 10498.607    6829.314 0.34950288    25
## 26 10800.558    6927.115 0.35863358    26
## 27 10114.764    7026.186 0.30535342    27
## 28  8570.413    7126.775 0.16844444    28
## 29  8711.170    7228.726 0.17017740    29
## 30  7532.565    7332.195 0.02660049    30
```

```
mean(prediction_table$`%Change`)
```

```
## [1] 0.3097158
```

```
ggplot() + geom_line(aes(x = Week, y = Prediction, color = "steelblue"),
  prediction_table) + geom_line(aes(x = Week, y = Actual, color = "red"),
  prediction_table) + ylim(0, 20000) + ylab("Bitcoin Price") +
  scale_color_manual(labels = c("Actual", "Prediction"), values = c("steelblue",
    "red")) + theme_classic() + ggtitle("Actual VS Predicted values")
```



## Add Google Trend As a Predictor

```
tsdata <- ts(training)

## Warning in data.matrix(data): NAs introduced by coercion


## Warning in data.matrix(data): NAs introduced by coercion

fit_trend <- auto.arima(training$log_price, xreg = training$google_trend)
fit_trend  #AIC=-405.75  AICc=-405.52  BIC=-387.97

## Series: training$log_price
## Regression with ARIMA(2,1,1) errors
##
## Coefficients:
##          ar1      ar2      ma1      drift      xreg
##        -0.0393  0.1705  0.2730  0.0142  0.0166
## s.e.      0.7104  0.1895  0.7203  0.0088  0.0046
##
## sigma^2 estimated as 0.008456: log likelihood=224.04
## AIC=-436.08  AICc=-435.7  BIC=-415.48
# Lower AIC, AICc and BIC compare with model without google
# trend

tsdisplay(residuals(fit_trend), lag.max = 30, main = "(1,1,1) Model Residuals")
```


**(1,1,1) Model Residuals**



Check Residuals for model validity


```
res1 <- fit_trend$residuals  
Acf(res1, main = "ACF for Residuals")
```

### ACF for Residuals



```
Pacf(res1, main = "PACF for Residuals")
```

## PACF for Residuals



```
# Based on ACF and PACF, residuals are randomly distributed.
```

```
Box.test(res1, type = "Ljung", lag = 1)
```

```
##
```

```
## Box-Ljung test
```

```
##
```

```
## data: res1
```

```
## X-squared = 0.014416, df = 1, p-value = 0.9044
```

```
# P-value is high so we fail to reject H0: The residuals are  
# independent. Box-Ljung test confirms our results.
```

## Google Trend ARIMA Model Forecast

```
# Forecast
```

```
tsdata_trend <- ts(testing$google_trend)
```

```
fc.c2 <- forecast(tsdata_trend, h = 30)
```


```
newxreg <- as.matrix(fc.c2$mean)
```

```
fcast_trend <- forecast(fit_trend, xreg = newxreg)
```

```
plot(fcast_trend)
```




## Forecasts from Regression with ARIMA(2,1,1) errors



```
# Take the exponent
fcast_trend$mean <- antilog(fcast_trend$mean)
fcast_trend$lower <- antilog(fcast_trend$lower)
fcast_trend$upper <- antilog(fcast_trend$upper)
fcast_trend$x <- antilog(fcast_trend$x)

plot(fcast_trend)
```

## Forecasts from Regression with ARIMA(2,1,1) errors



## Compare prediction and actual bitcoin price for 30 weeks


```
actual <- testing$bitcoin_price
pred1 <- fcast_trend$mean[1:30]
error1 <- abs(actual - pred1)/actual
prediction_table1 <- cbind(actual, pred1, error1)
colnames(prediction_table1) <- c("Actual", "Prediction", "%Change")
prediction_table1 <- as.data.frame(prediction_table1)
prediction_table1$Week <- seq(1:30)
prediction_table1
```

##	Actual	Prediction	%Change	Week
## 1	4448.186	4846.434	0.08953034	1
## 2	3877.361	4960.340	0.27930843	2
## 3	3836.227	5045.311	0.31517517	3
## 4	4104.402	5124.725	0.24859233	4
## 5	4349.547	5200.129	0.19555662	5
## 6	5163.845	5275.623	0.02164628	6
## 7	5724.994	5351.330	0.06526878	7
## 8	5764.160	5427.980	0.05832249	8
## 9	6814.830	5505.579	0.19211786	9
## 10	6936.828	5584.268	0.19498243	10
## 11	7254.876	5664.057	0.21927582	11
## 12	8305.702	5744.983	0.30830857	12
## 13	10326.294	5827.061	0.43570647	13
## 14	13732.345	5910.311	0.56960655	14
## 15	17376.360	5994.750	0.65500541	15
## 16	16223.990	6080.396	0.62522191	16
## 17	14485.893	6167.265	0.57425721	17
## 18	15325.235	6255.375	0.59182521	18
## 19	14512.732	6344.743	0.56281538	19
## 20	12177.336	6435.389	0.47152739	20
## 21	11282.490	6527.329	0.42146376	21
## 22	9788.082	6620.583	0.32360768	22
## 23	8014.081	6715.170	0.16207860	23
## 24	9670.893	6811.107	0.29571060	24
## 25	10498.607	6908.416	0.34196836	25
## 26	10800.558	7007.114	0.35122665	26
## 27	10114.764	7107.223	0.29734170	27
## 28	8570.413	7208.761	0.15887821	28
## 29	8711.170	7311.751	0.16064654	29
## 30	7532.565	7416.212	0.01544676	30

```
mean(prediction_table1$`%Change`)
```

```
## [1] 0.3067473
```

```
ggplot() + geom_line(aes(x = Week, y = Prediction, color = "steelblue"),
  prediction_table1) + geom_line(aes(x = Week, y = Actual,
  color = "red"), prediction_table1) + ylim(0, 20000) + ylab("Bitcoin Price") +
  scale_color_manual(labels = c("Actual", "Prediction"), values = c("steelblue",
  "red")) + theme_classic() + ggtitle("Actual VS Predicted values")
```



## Accuracy of all ARIMA Model

```
# Accuracy
accuracy(prediction_table$Actual, prediction_table$Prediction)

##              ME      RMSE      MAE      MPE      MAPE
## Test set -3183.876 4757.224 3468.105 -50.97236 56.66672

accuracy(prediction_table1$Actual, prediction_table1$Prediction)

##              ME      RMSE      MAE      MPE      MAPE
## Test set -3111.502 4708.59 3423.035 -49.1853 55.34531
```